

Avocado Data Analysis

Business Understanding

The aim of this project is to answer the following four questions: 1. Which region are the lowest and highest prices of Avocado? 2. What is the highest region of avocado production? 3. What is the average avocado prices in each year? 4. What is the average avocado volume in each year?

Data Understanding

The [Avocado dataset](#) was been used in this project.

This dataset contains 13 columns: 1. Date - The date of the observation 2. AveragePrice: the average price of a single avocado 3. Total Volume: Total number of avocados sold 4. Total Bags: Total number o bags 5. Small Bags: Total number of Small bags 6. Large Bags: Total number of Large bags 7. XLarge Bags: Total number of XLarge bags 8. type: conventional or organic 9. year: the year 10. region: the city or region of the observation 11. 4046: Total number of avocados with PLU 4046 sold 12. 4225: Total number of avocados with PLU 4225 sold 13. 4770: Total number of avocados with PLU 4770 sold

Import necessary libraries

```
In [72]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import r2_score
from sklearn.pipeline import Pipeline
```

Data preparation

Load data

```
In [73]: df = pd.read_csv(r"c:\Users\Hanshu\Desktop\excel data_ML\avocado.csv")
df
```

Out[73]:

	Unnamed: 0	Date	AveragePrice	Total Volume	4046	4225	4770	Total Bags
0	0	2015-12-27	1.33	64236.62	1036.74	54454.85	48.16	8696.87
1	1	2015-12-20	1.35	54876.98	674.28	44638.81	58.33	9505.56
2	2	2015-12-13	0.93	118220.22	794.70	109149.67	130.50	8145.35
3	3	2015-12-06	1.08	78992.15	1132.00	71976.41	72.58	5811.16
4	4	2015-11-29	1.28	51039.60	941.48	43838.39	75.78	6183.95
...
18244	7	2018-02-04	1.63	17074.83	2046.96	1529.20	0.00	13498.67
18245	8	2018-01-28	1.71	13888.04	1191.70	3431.50	0.00	9264.84
18246	9	2018-01-21	1.87	13766.76	1191.92	2452.79	727.94	9394.11
18247	10	2018-01-14	1.93	16205.22	1527.63	2981.04	727.01	10969.54
18248	11	2018-01-07	1.62	17489.58	2894.77	2356.13	224.53	12014.15

18249 rows × 14 columns



Explore the data

```
In [74]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 18249 entries, 0 to 18248
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0            18249 non-null  int64
1   Date                  18249 non-null  object
2   AveragePrice          18249 non-null  float64
3   Total Volume         18249 non-null  float64
4   4046                  18249 non-null  float64
5   4225                  18249 non-null  float64
6   4770                  18249 non-null  float64
7   Total Bags            18249 non-null  float64
8   Small Bags            18249 non-null  float64
9   Large Bags            18249 non-null  float64
10  XLarge Bags           18249 non-null  float64
11  type                  18249 non-null  object
12  year                  18249 non-null  int64
13  region                18249 non-null  object
dtypes: float64(9), int64(2), object(3)
memory usage: 1.9+ MB

```

In [75]: `df.head()`

Out[75]:

	Unnamed: 0	Date	AveragePrice	Total Volume	4046	4225	4770	Total Bags	Small Bags
0	0	2015-12-27	1.33	64236.62	1036.74	54454.85	48.16	8696.87	8603
1	1	2015-12-20	1.35	54876.98	674.28	44638.81	58.33	9505.56	9408
2	2	2015-12-13	0.93	118220.22	794.70	109149.67	130.50	8145.35	8042
3	3	2015-12-06	1.08	78992.15	1132.00	71976.41	72.58	5811.16	5677
4	4	2015-11-29	1.28	51039.60	941.48	43838.39	75.78	6183.95	5986

Missing value checking

In [76]: `df.isnull().sum()`

```
Out[76]: Unnamed: 0      0
         Date          0
         AveragePrice  0
         Total Volume  0
         4046          0
         4225          0
         4770          0
         Total Bags    0
         Small Bags    0
         Large Bags    0
         XLarge Bags   0
         type          0
         year          0
         region        0
         dtype: int64
```

Dropping unnecessary columns

```
In [77]: df = df.drop(['Unnamed: 0', '4046', '4225', '4770', 'Date'], axis=1)
```

```
In [78]: df.head()
```

```
Out[78]:
```

	AveragePrice	Total Volume	Total Bags	Small Bags	Large Bags	XLarge Bags	type	year	region
0	1.33	64236.62	8696.87	8603.62	93.25	0.0	conventional	2015	Alban
1	1.35	54876.98	9505.56	9408.07	97.49	0.0	conventional	2015	Alban
2	0.93	118220.22	8145.35	8042.21	103.14	0.0	conventional	2015	Alban
3	1.08	78992.15	5811.16	5677.40	133.76	0.0	conventional	2015	Alban
4	1.28	51039.60	6183.95	5986.26	197.69	0.0	conventional	2015	Alban

Answering questions

```
In [79]: def get_average(df, column):
         """
         Description: This function to return the average value of the column

         Arguments:
             df: the DataFrame.
             column: the selected column.
         Returns:
             column's average
         """
         return sum(df[column])/len(df)
```

```
In [80]: def get_avarge_between_two_columns(df, column1, column2):
         """
         Description: This function calculate the average between two columns in the

         Arguments:
             df: the DataFrame.
             column1: the first column.
```

```

        column2:the scnd column.
Returns:
    Sorted data for relation between column1 and column2
    """

List = list(df[column1].unique())
average=[]
for i in List:
    x = df[df[column1]==i]
    column1_average= get_average(x,column2)
    average.append(column1_average)

df_column1_column2 = pd.DataFrame({'column1':List,'column2':average})
column1_column2_sorted_index = df_column1_column2.column2.sort_values(ascending=True)
column1_column2_sorted_data=df_column1_column2.reindex(column1_column2_sorted_index)

return column1_column2_sorted_data

```

```

In [81]: def plot(data,xlabel,ylabel):
    """
    Description: This function to draw a barplot

    Arguments:
        data: the DataFrame.
        xlabel: the label of the first column.
        ylabel: the label of the second column.
    Returns:
        None
    """

    plt.figure(figsize=(15,5))
    ax = sns.barplot(data=data,x=data.columns[0],y=data.columns[1])
    plt.xticks(rotation=90)
    plt.xlabel(xlabel)
    plt.ylabel(ylabel)
    plt.title((' Average'+ylabel+' of Avocado According to '+xlabel))

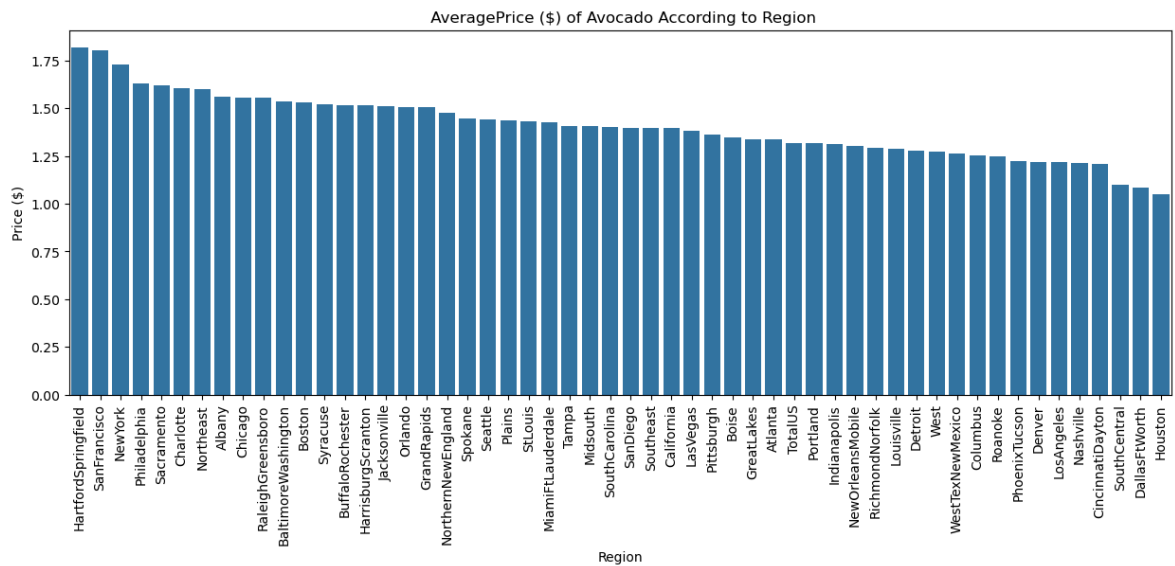
```

Which region are the lowest and highest prices of Avocado?

```

In [82]: data1 = get_avarge_between_two_columns(df,'region','AveragePrice')
plot(data1,'Region','Price ($)')

```



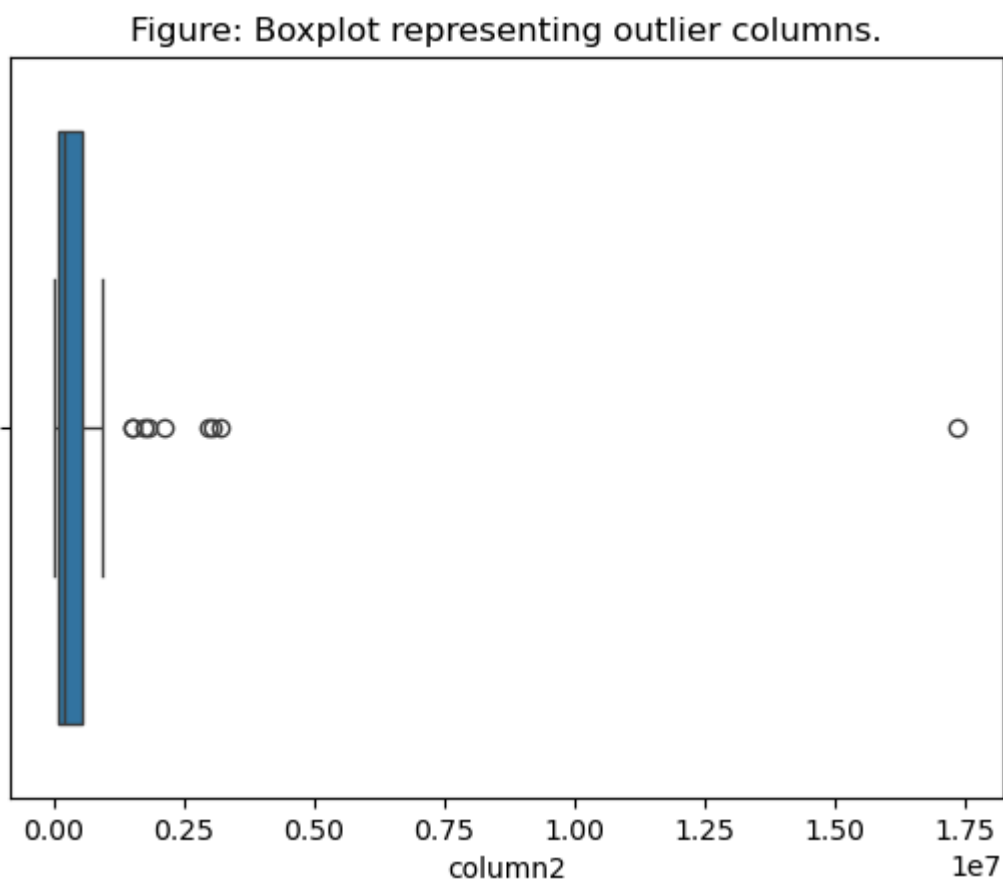
```
In [83]: print(data1['column1'].iloc[-1], " is the region producing avocado with the lowe
Houston is the region producing avocado with the lowest price.
```

What is the highest region of avocado production?

Checking if there are outlier values or not.

```
In [84]: data2 = get_avarge_between_two_columns(df,'region', 'Total Volume')
sns.boxplot(x=data2.column2).set_title("Figure: Boxplot representing outlier col
```

```
Out[84]: Text(0.5, 1.0, 'Figure: Boxplot representing outlier columns.')
```




```
In [85]: outlier_region = data2[data2.column2>10000000]
print(outlier_region['column1'].iloc[-1], "is outlier value")
```

TotalUS is outlier value

Remove the outlier values

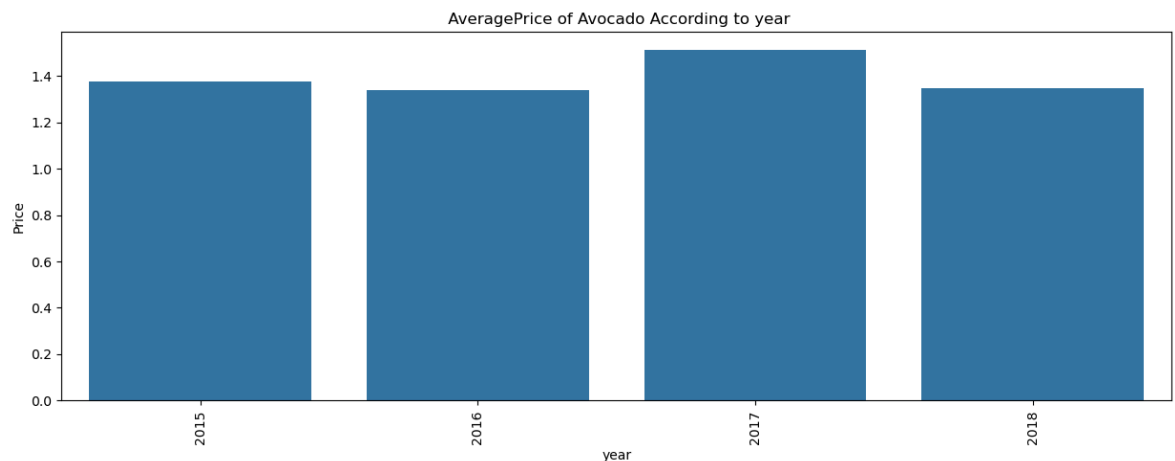
```
In [86]: outlier_region.index
data2 = data2.drop(outlier_region.index,axis=0)
```

```
In [87]: plot(data2,'Region','Volume')
```




What is the average avocado prices in each year?

```
In [88]: data3 = get_avarge_between_two_columns(df,'year','AveragePrice')
plot(data3,'year','Price')
```



What is the average avocado volume in each year?

```
In [89]: data4 = get_avarge_between_two_columns(df,'year','Total Volume')
plot(data4,'year','Volume')
```



Data Modeling

We built the regression model by using [Linear regression from sklearn](#) to predict the avocado price.

Changing some column types to categories

```
In [90]: df['region'] = df['region'].astype('category')
df['region'] = df['region'].cat.codes

df['type'] = df['type'].astype('category')
df['type'] = df['type'].cat.codes
```

```
In [91]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 18249 entries, 0 to 18248
Data columns (total 9 columns):
#   Column      Non-Null Count  Dtype
---  -
0   AveragePrice 18249 non-null  float64
1   Total Volume 18249 non-null  float64
2   Total Bags   18249 non-null  float64
3   Small Bags   18249 non-null  float64
4   Large Bags   18249 non-null  float64
5   XLarge Bags  18249 non-null  float64
6   type         18249 non-null  int8
7   year         18249 non-null  int64
8   region       18249 non-null  int8
dtypes: float64(6), int64(1), int8(2)
memory usage: 1.0 MB
```

```
In [92]: df.head()
```


Out[92]:

	AveragePrice	Total Volume	Total Bags	Small Bags	Large Bags	XLarge Bags	type	year	region
0	1.33	64236.62	8696.87	8603.62	93.25	0.0	0	2015	0
1	1.35	54876.98	9505.56	9408.07	97.49	0.0	0	2015	0
2	0.93	118220.22	8145.35	8042.21	103.14	0.0	0	2015	0
3	1.08	78992.15	5811.16	5677.40	133.76	0.0	0	2015	0
4	1.28	51039.60	6183.95	5986.26	197.69	0.0	0	2015	0

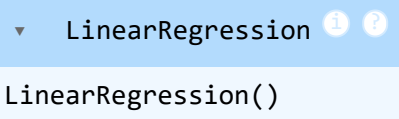
```
In [93]: # split data into X and y
X = df.drop(['AveragePrice'],axis=1)
y = df['AveragePrice']

# split data into traing and testing dataset
X_train, X_test, y_train, y_test = train_test_split(X,
                                                    y,
                                                    test_size=0.3,
                                                    random_state=15)
```

```
In [98]: print('training set:',X_train.shape,' - ',y_train.shape[0],' samples')
print("testing set:",X_test.shape,' - ',y_test.shape[0],' samples')
```

training set: (12774, 8) - 12774 samples
testing set: (5475, 8) - 5475 samples

```
In [103... # bulid and fit the model
model = LinearRegression()
model.fit(X_train, y_train)
```

Out[103...  LinearRegression()

Evaluate the Results

```
In [104... # prediction and calculate the accuracy for the testing dataset
test_pre = model.predict(X_test)
test_score = r2_score(y_test,test_pre)
print("The accuracy of testing dataset ",test_score*100)
```

The accuracy of testing dataset 38.58074176447186

```
In [105... # prediction and calculate the accuracy for the testing dataset
train_pred = model.predict(X_train)
train_score = r2_score(y_train, train_pred)
print("The accuracy of ttraining dataset", train_score*100)
```

The accuracy of ttraining dataset 39.706860424100924

The model doesn't work well with this dataset, In order to the avocado prices were near together